

SP Programming assignment 4 Report

B03902086 李鈺昇

1. The data of 24 running results (seconds)

(The output is shown in Appendix A.)

Size = 100:

Seg. size	1	4	10	20	50	100
Real	0.013	0.004	0.003	0.003	0.005	0.003
User	0.003	0.001	0.001	0.001	0.001	0.001
Sys	0.011	0.003	0.002	0.002	0.001	0.002

Size = 10000:

Seg. size	100	400	1000	2000	5000	10000
Real	0.027	0.017	0.014	0.013	0.010	0.009
User	0.022	0.016	0.013	0.012	0.009	0.007
Sys	0.009	0.003	0.002	0.002	0.001	0.001

Size = 1000000:

Seg. size	10000	40000	100000	200000	500000	1000000
Real	1.591	1.248	1.143	0.947	0.761	0.679
User	1.787	1.425	1.293	1.068	0.832	0.664
Sys	0.032	0.017	0.013	0.010	0.009	0.005

Size = 10000000:

Seg. size	100000	400000	1000000	2000000	5000000	10000000
Real	17.301	15.872	12.379	10.781	7.716	7.395
User	18.414	16.659	13.698	11.796	8.534	7.108
Sys	0.377	0.363	0.262	0.188	0.076	0.064

2. My observation and discovery

If the input size is fixed (look at each individual table), then the trend is that as the segment size becomes larger, both user time and system time decrease, and hence so does real time. I think the main reason lies in the decrease of the number of segments when segment size becomes larger. Smaller number of segments means fewer times of merging two sorted lists, which is the main bottleneck of merge-sort. Although sorting with fewer segments seems less parallel, the sorting part is not what slows down the whole process. In fact, the time of merging dominates the time of sorting.

Next, if the ratio of the size of segments to the input size is fixed (look at each columns across the four tables), it seems that the real time is in proportion to the input size. This is reasonable without any question.

A question is here: why is the real time, in some cases, less than the sum of CPU time and system time? Soon I realized. Since we're running a multi-thread program, the CPU time is the sum of all the running time of the separate CPUs, while the real time may be just around the maximum of them. This is why CPU time can exceed real time in this context.

3. Experiment 1: The growth of time as segment size becomes smaller

(The output is shown in Appendix B.)

When the input size is 10000000, the minimal segment size we are asked to run is 100000, which results in the longest total real time. I wonder how the real time changes as the segment size becomes smaller. I fix input size at 10000000, and let the segment size be 10000000, 1000000, 100000, ..., each term being the previous term divided by ten. The result is as follows:

Seg. size	10000000	1000000	100000	10000	1000	100
Real	7.591	11.178	17.652	21.086	18.639	159.473
User	7.216	12.898	19.127	23.454	19.536	13.334
Sys	0.076	0.137	0.326	0.420	0.698	202.654

*Segment sizes under 10 are not listed, because they took too much time (> 30mins) without results.

Clearly, cases where segment sizes are greater than or equal to 1000 are not too far from one

another. However, the case where segment size is 100 took explosive system time (user time is still fine). Two or three minutes are much longer than other cases, so the case of 10 must be a lot longer than the listed cases*. Thus 100 might be near the critical value where the running time grows suddenly faster.

4. Experiments 2: The comparison of “long” and “int”

(The output is shown in Appendix C.)

At first, I encounter the problem of conversion, or type casting, between `int` and `void*`. Converting between `int` and `void*` generated warnings, which I disliked. Then I found that this problem could be resolved if `int` was replaced with `long`. Then everything really works fine. But I wonder how the performance of them differ, so I changed `long` back to `int`, ignoring the warnings, and then re-compile, re-run again.

At first I guessed that `int` version would be faster than `long` version. Surprisingly, the `long` version is faster overall, in many cases! Then I thought the reason behind this might be the time spent on converting between `int` and `void*`, which may take longer time than between `long` and `void*`.

Appendix A: Original output

```
size: 100, seg_size: 1
```

```
real 0m0.013s
user 0m0.003s
sys 0m0.011s
```

```
=====
```

```
size: 100, seg_size: 4
```

```
real 0m0.004s
user 0m0.001s
sys 0m0.003s
```

```
=====
```

```
size: 100, seg_size: 10
```

```
real 0m0.003s
user 0m0.001s
sys 0m0.002s
```

```
=====
```

```
size: 100, seg_size: 20
```

```
real 0m0.003s
user 0m0.001s
```

```
sys 0m0.002s
=====
size: 100, seg_size: 50

real 0m0.005s
user 0m0.001s
sys 0m0.001s
=====
size: 100, seg_size: 100

real 0m0.003s
user 0m0.001s
sys 0m0.002s
=====
size: 10000, seg_size: 100

real 0m0.027s
user 0m0.022s
sys 0m0.009s
=====
size: 10000, seg_size: 400

real 0m0.017s
user 0m0.016s
sys 0m0.003s
=====
size: 10000, seg_size: 1000

real 0m0.014s
user 0m0.013s
sys 0m0.002s
=====
size: 10000, seg_size: 2000

real 0m0.013s
user 0m0.012s
sys 0m0.002s
=====
size: 10000, seg_size: 5000

real 0m0.010s
user 0m0.009s
sys 0m0.001s
=====
size: 10000, seg_size: 10000

real 0m0.009s
user 0m0.007s
sys 0m0.001s
=====
size: 1000000, seg_size: 10000

real 0m1.591s
user 0m1.787s
sys 0m0.032s
=====
size: 1000000, seg_size: 40000

real 0m1.248s
user 0m1.425s
```

```
sys 0m0.017s
=====
size: 1000000, seg_size: 100000

real 0m1.143s
user 0m1.293s
sys 0m0.013s
=====
size: 1000000, seg_size: 200000

real 0m0.947s
user 0m1.068s
sys 0m0.010s
=====
size: 1000000, seg_size: 500000

real 0m0.761s
user 0m0.832s
sys 0m0.009s
=====
size: 1000000, seg_size: 1000000

real 0m0.679s
user 0m0.664s
sys 0m0.005s
=====
size: 10000000, seg_size: 100000

real 0m17.301s
user 0m18.414s
sys 0m0.377s
=====
size: 10000000, seg_size: 400000

real 0m15.872s
user 0m16.659s
sys 0m0.363s
=====
size: 10000000, seg_size: 1000000

real 0m12.379s
user 0m13.698s
sys 0m0.262s
=====
size: 10000000, seg_size: 2000000

real 0m10.781s
user 0m11.796s
sys 0m0.188s
=====
size: 10000000, seg_size: 5000000

real 0m7.716s
user 0m8.534s
sys 0m0.076s
=====
size: 10000000, seg_size: 10000000

real 0m7.395s
user 0m7.108s
```

```
sys 0m0.064s
=====
```

Appendix B: Output of experiment 1

(The last case didn't complete, and was terminated.)

```
size: 10000000, seg_size: 10000000

real 0m7.591s
user 0m7.216s
sys 0m0.076s
=====
size: 10000000, seg_size: 1000000

real 0m11.178s
user 0m12.898s
sys 0m0.137s
=====
size: 10000000, seg_size: 100000

real 0m17.652s
user 0m19.127s
sys 0m0.326s
=====
size: 10000000, seg_size: 10000

real 0m21.086s
user 0m23.454s
sys 0m0.420s
=====
size: 10000000, seg_size: 1000

real 0m18.639s
user 0m19.536s
sys 0m0.698s
=====
size: 10000000, seg_size: 100

real 2m39.473s
user 0m13.334s
sys 3m22.654s
=====
size: 10000000, seg_size: 10
```

Appendix C: Output of experiment 2

```
size: 100, seg_size: 1

real 0m0.010s
user 0m0.003s
sys 0m0.010s
=====
```

```
size: 100, seg_size: 4

real 0m0.006s
user 0m0.001s
sys 0m0.004s
=====
size: 100, seg_size: 10

real 0m0.004s
user 0m0.001s
sys 0m0.002s
=====
size: 100, seg_size: 20

real 0m0.012s
user 0m0.001s
sys 0m0.002s
=====
size: 100, seg_size: 50

real 0m0.005s
user 0m0.001s
sys 0m0.002s
=====
size: 100, seg_size: 100

real 0m0.003s
user 0m0.001s
sys 0m0.001s
=====
size: 10000, seg_size: 100

real 0m0.026s
user 0m0.022s
sys 0m0.009s
=====
size: 10000, seg_size: 400

real 0m0.018s
user 0m0.017s
sys 0m0.003s
=====
size: 10000, seg_size: 1000

real 0m0.017s
user 0m0.014s
sys 0m0.003s
=====
size: 10000, seg_size: 2000

real 0m0.013s
user 0m0.011s
sys 0m0.002s
=====
size: 10000, seg_size: 5000

real 0m0.010s
user 0m0.009s
sys 0m0.001s
=====
```

```
size: 10000, seg_size: 10000

real 0m0.009s
user 0m0.007s
sys 0m0.001s
=====
size: 1000000, seg_size: 10000

real 0m1.648s
user 0m1.838s
sys 0m0.036s
=====
size: 1000000, seg_size: 40000

real 0m1.378s
user 0m1.570s
sys 0m0.021s
=====
size: 1000000, seg_size: 100000

real 0m1.179s
user 0m1.316s
sys 0m0.018s
=====
size: 1000000, seg_size: 200000

real 0m1.014s
user 0m1.126s
sys 0m0.015s
=====
size: 1000000, seg_size: 500000

real 0m0.785s
user 0m0.857s
sys 0m0.010s
=====
size: 1000000, seg_size: 1000000

real 0m0.703s
user 0m0.691s
sys 0m0.008s
=====
size: 10000000, seg_size: 100000

real 0m18.917s
user 0m20.368s
sys 0m0.430s
=====
size: 10000000, seg_size: 400000

real 0m16.482s
user 0m16.893s
sys 0m0.443s
=====
size: 10000000, seg_size: 1000000

real 0m12.898s
user 0m14.062s
sys 0m0.289s
=====
```



```

size: 10000000, seg_size: 2000000

real 0m11.529s
user 0m12.234s
sys 0m0.272s
=====
size: 10000000, seg_size: 5000000

real 0m10.083s
user 0m9.877s
sys 0m0.258s
=====
size: 10000000, seg_size: 10000000

real 0m9.613s
user 0m8.044s
sys 0m0.213s
=====

```

Appendix D: Shell script to generate Appendix A, C

(in10, ..., in10000000 are the random input files.)

```

#!/bin/bash

size[0]=100
size[1]=10000
size[2]=1000000
size[3]=10000000

nseg[0]=100
nseg[1]=25
nseg[2]=10
nseg[3]=5
nseg[4]=2
nseg[5]=1

for i in {0..3}
do
    for j in {0..5}
    do
        echo "size: ${size[$i]}, seg_size: $(( ${size[$i]} / ${nseg[$j]} ))"
        time ./merger $(( ${size[$i]} / ${nseg[$j]} )) < in${size[$i]} > /dev/null
        echo "===== "
    done
done

```

Appendix E: Shell script to generate Appendix B

```

#!/bin/bash

size[0]=10000000
size[1]=1000000
size[2]=100000
size[3]=10000

```

```
size[4]=1000
size[5]=100
size[6]=10
size[7]=1

for i in {0..7}
do
    echo "size: 10000000, seg_size: ${size[$i]}"
    time ./merger ${size[$i]} < in10000000 > /dev/null
    echo "===== "
done
```